

# **Concurrency Control: 2PL**

Dr. L.M. Jenila Livingston  
VIT Chennai

# *Lock Management*

- A lock is a mechanism to **control concurrent access** to a data item
- Lock requests are made to **concurrency-control manager** or **lock manager**. Transaction can proceed only after request is granted.
- **Lock table** entry:
  - **Number of transactions** currently holding a lock
  - **Type of lock** held (shared or exclusive)
  - **Pointer to queue** of lock requests

# 3 Types of Locks

- We allow transactions to **lock** objects.

**Shared lock (S):** Data item can **only be read**. S-lock is requested using **lock-S** instruction.

**Exclusive lock (X):** Data item can be **both read** as well as **write**. X-lock is requested using **lock-X** instruction.

# Lock-Based Protocols (Cont.)

## ■ Lock-compatibility matrix


	S	X
S	true	false
X	false	false

- **Any number of transactions can hold shared locks on an item,**
  - but if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.

# The Two-Phase Locking (2PL) Protocol

- This is a protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
  - transaction may obtain locks
  - transaction may not release locks
- Phase 2: Shrinking Phase
  - transaction may release locks
  - transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e. the point where a transaction acquired its final lock).

## 2 PL



# The Two-Phase Locking Protocol (Cont.)

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules
- Two-phase locking *does not* ensure freedom from deadlocks
- **Cascading roll-back is possible** under two-phase locking. To avoid this, follow a modified protocol called **strict two-phase locking**. Here a transaction must hold all its **exclusive locks** till it commits/aborts.
- **Rigorous two-phase locking** is even stricter: here **all locks** are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.
- .

# Lock Conversions

- Two-phase locking with lock conversions:
  - First Phase:
    - can acquire a lock-S on item
    - can acquire a lock-X on item
    - can convert a lock-S to a lock-X (upgrade)
  - Second Phase:
    - can release a lock-S
    - can release a lock-X
    - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.



# 2 PL: Automatic Acquisition of Locks

## - read

- A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- The operation **read**( $D$ ) is processed as:

```
if  $T_i$  has a lock on  $D$  // lock-X OR lock-S
then
    read( $D$ )
else begin
    if necessary wait until no other
        transaction has a lock-X on  $D$ 
    grant  $T_i$  a lock-S on  $D$ ;
    read( $D$ )
end
```

# 2PL: Automatic Acquisition of Locks

## - write

- **write( $D$ )** is processed as:

if  $T_i$  has a **lock-X** on  $D$

**then**

    write( $D$ )

**else begin**

    if necessary wait until no other trans. has any lock on  $D$ ,

    if  $T_i$  has a **lock-S** on  $D$

**then**

**upgrade** lock on  $D$  to **lock-X**

**else**

        grant  $T_i$  a **lock-X** on  $D$

        write( $D$ )

**end;**

- **All locks are released after commit or abort**

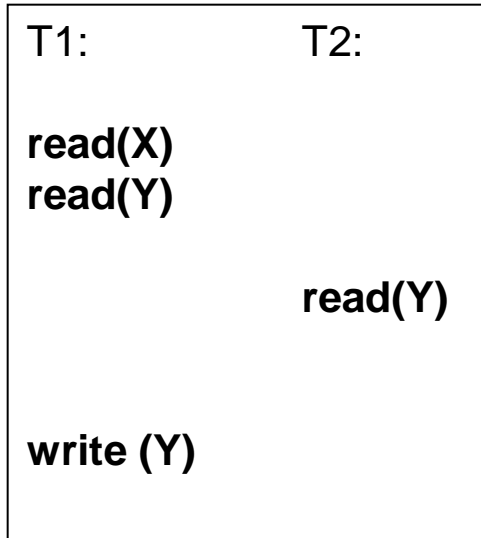
# Lock-Based Protocols (Cont.)

- **Example** of a transaction performing locking:

$T_2$ : **begin**  
    **lock-S**(A);  
    **read** (A);  
    **unlock**(A);  
    **lock-S**(B);  
    **read** (B);  
    **unlock**(B);  
    **display**(A+B)  
    **commit**;

<p>T2:      <b>read</b>(A)     <b>read</b>(B)     <b>display</b>(A+B)</p>
---

# Lock-Based Protocols – 2PL



T1:

begin  
lock-S(X)  
read(X)  
lock-S(Y)  
read(Y)

upgrade(Y)  
write (Y)  
unlock(X)  
unlock(Y)  
commit

T2:

begin  
lock-S(Y)  
read(Y)  
unlock(Y)  
commit

# Pitfalls of Lock-Based Protocols

- Consider the partial schedule

$T_3$	$T_4$
lock-x ( $B$ )	
read ( $B$ )	
$B := B - 50$	
write ( $B$ )	
	lock-s ( $A$ )
	read ( $A$ )
	lock-s ( $B$ )
lock-x ( $A$ )	

- Neither  $T_3$  nor  $T_4$  can make progress — executing **lock-S( $B$ )** causes  $T_4$  to wait for  $T_3$  to release its lock on  $B$ , while executing **lock-X( $A$ )** causes  $T_3$  to wait for  $T_4$  to release its lock on  $A$ .
- Such a situation is called a **deadlock**.
  - To handle a deadlock one of  $T_3$  or  $T_4$  must be rolled back and its locks released.


# Pitfalls of Lock-Based Protocols (Cont.)

- **Starvation** is also possible if concurrency control manager is badly designed. For example:
  - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
  - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.

# Implementation of Locking

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The **lock manager maintains a data-structure** called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked

# Lock Table



- Lock table records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
  - lock manager may keep a list of locks held by each transaction, to implement this efficiently



# Timestamp-Based Protocols

1. Use the value of the system clock as the timestamp; that is, a transaction's timestamp is equal to the value of the clock when the transaction enters the system.
2. Use a logical counter that is incremented after a new timestamp has been assigned; that is, a transaction's timestamp is equal to the value of the counter when the transaction enters the system.

# Timestamp-based Protocols

- Suppose there are there transactions T1, T2, and T3.
- T1 has entered the system at time 0010
- T2 has entered the system at 0020
- T3 has entered the system at 0030
- Priority will be given to transaction T1, then transaction T2 and lastly Transaction T3.

**Thank You!**

# **Transactions – Concurrent Executions**

Dr. L.M. Jenila Livingston  
VIT Chennai

# Transactions

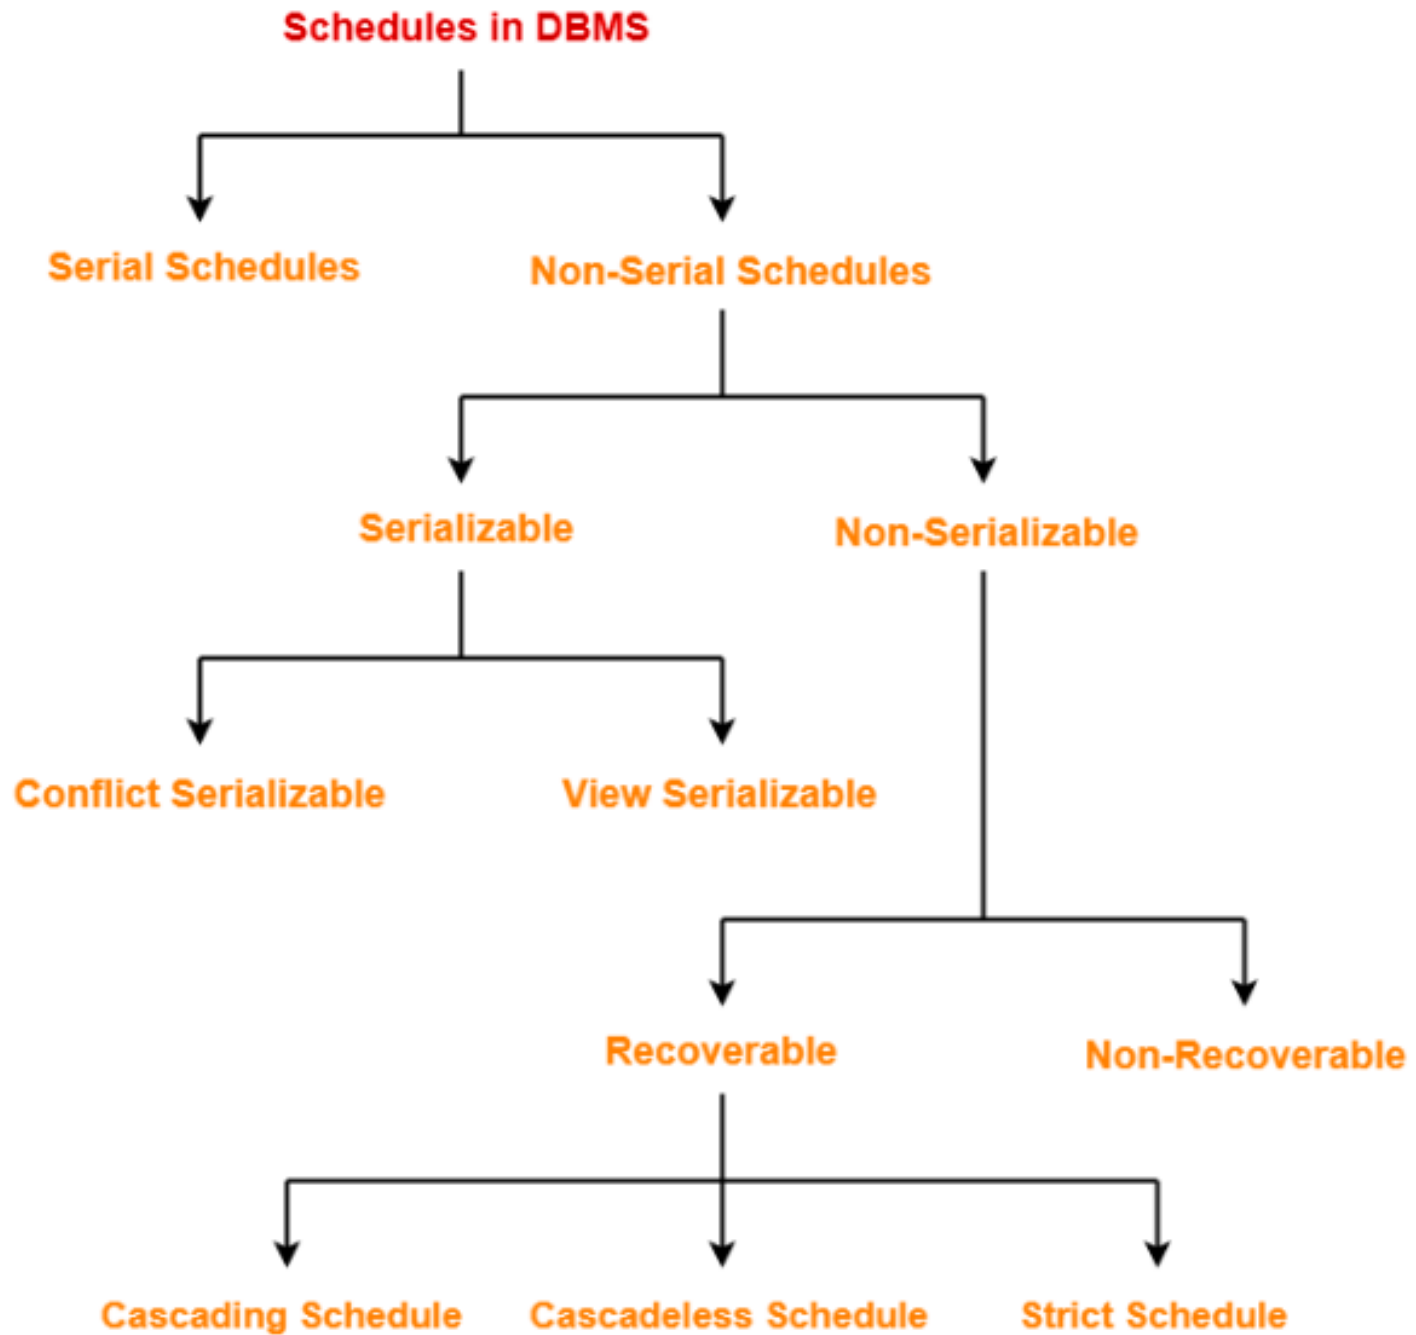
- Concurrent Executions
- Serializability
- Conflict Serializable
- Testing for Serializability
- View Serializable
- Recoverability
- Transaction Definition in SQL
- Recoverable Schedules

# Concurrent Executions

- **Multiple transactions are allowed to run concurrently** in the system.  
Advantages are:
  - **increased processor and disk utilization**, leading to better transaction *throughput*
    - ▶ E.g. one transaction can be using the CPU while another is reading from or writing to the disk
  - **reduced average response time** for transactions: short transactions need not wait behind long ones.
- **Concurrency control schemes** – mechanisms to achieve isolation
  - that is, to **control the interaction among the concurrent transactions** in order to prevent them from destroying the consistency of the database

# Schedules

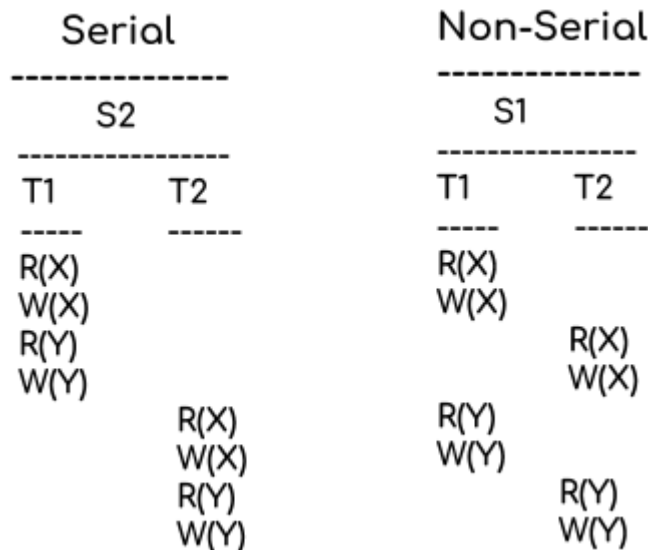
- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - a schedule for a set of transactions must consist of all instructions of those transactions
  - must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a **commit** instructions as the last statement
  - by default transaction assumed to **execute commit instruction as its last step**
- A transaction that fails to successfully complete its execution will have an **abort instruction** as the last statement





# Serial Vs Serializable

Serial Schedules	Serializable Schedules
<p><b>No concurrency</b> is allowed. Thus, all the transactions necessarily execute serially one after the other.</p>	<p><b>Concurrency</b> is allowed. Thus, multiple transactions can execute concurrently.</p>
Serial schedules lead to less resource utilization and CPU throughput.	Serializable schedules improve both resource utilization and CPU throughput.
Serial Schedules are less efficient as compared to serializable schedules.	Serializable Schedules are always better than serial schedules.



# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- A **serial schedule** in which  $T_1$  is followed by  $T_2$  :

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Let  $A=150$ ,  $B=50$   
 $A=150-50=100$   
 $B=50+50=100$

$temp=100*0.1=10$   
 $A=100-10=90$   
 $B=100+10=110$

Sum( $A+B$ )=200

# Schedule 2

- A **serial schedule** where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Let  $A=150$ ,  $B=50$   
 $Temp=150*0.1=15$   
 $A=150-15=135$   
 $B=50+15=65$

$A=135-50=85$   
 $B=65+50=115$

$Sum(A+B)=200$

# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1.

$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Let A=150, B=50  
 $A = 150 - 50 = 100$

$temp = 100 * 0.1 = 10$   
 $A = 100 - 10 = \mathbf{90}$

$B = 50 + 50 = 100$

$B = 100 + 10 = \mathbf{110}$

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

$Sum(A+B) = 200$

# Schedule 4


- The following concurrent schedule does not preserve the value of  $(A + B)$ .

$T_1$	$T_2$
read (A) $A := A - 50$	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	
	$B := B + temp$ write (B) commit

Let  $A=150$ ,  $B=50$   
 $A=150-50$   
  
 $temp=150*0.1=15$   
 $A=150-15=135$   
  
 **$A=100$**   
 **$B=50+50=100$**   
  
 $B=50+15=65$

# Serializability


- **Basic Assumption** – Each transaction preserves database consistency.
- Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is **equivalent** to a serial schedule.
- Different forms of schedule equivalence give rise to the notions of:
  1. **view serializability**
  2. **conflict serializability**



# Serializability

Serializable: A schedule is serializable if its effects on the database are the equivalent to some serial schedule.

Hard to ensure; more conservative approaches are used in practice

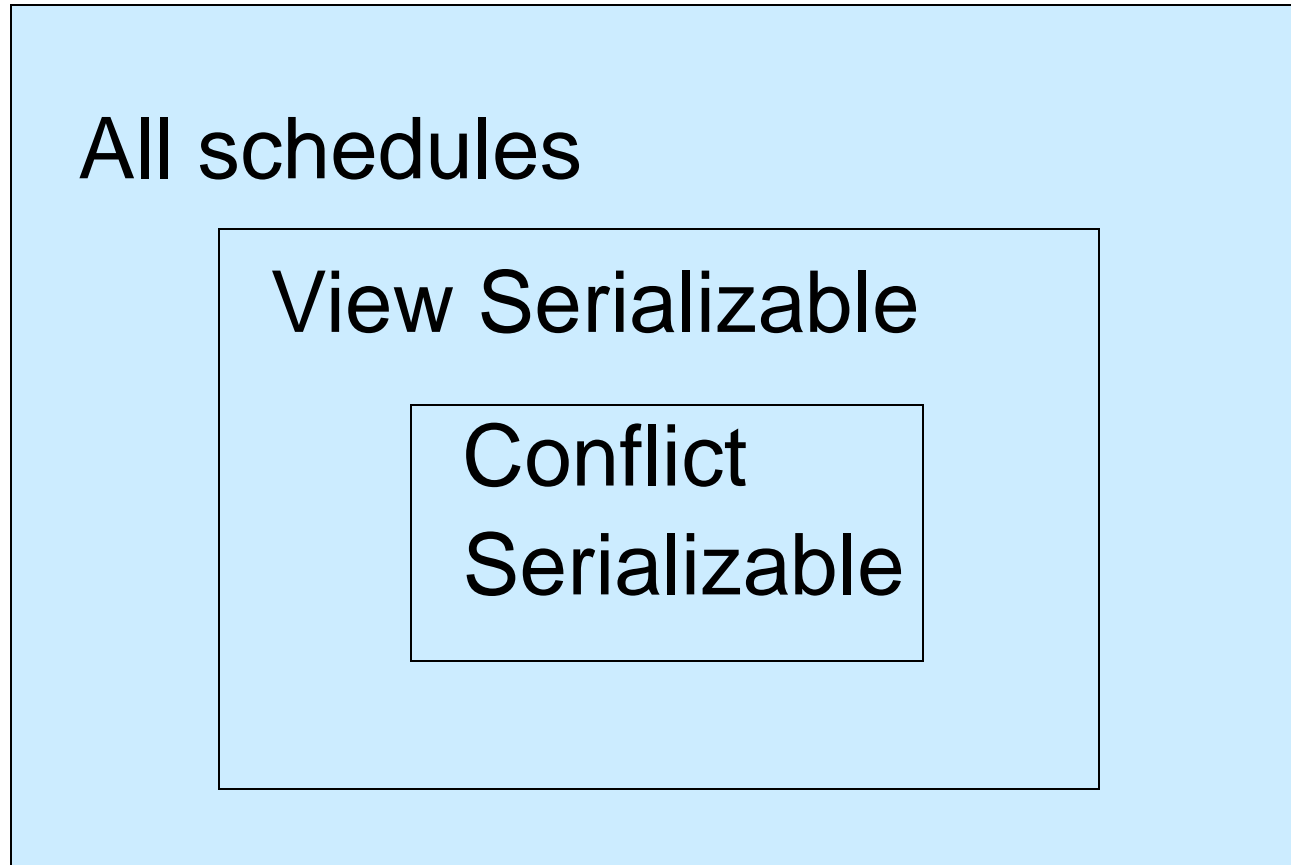


# ***Simplified view of transactions***

- We **ignore operations** other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- Our simplified schedules consist of only **read** and **write** instructions.



# Diagram



# Conflicts

## ■ Definition

- is a pair of consecutive actions in a schedule such that, if their **order is interchanged**, then the **behavior** of at least one of the transactions involved **can change**.

# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ .
  1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict.
  2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict.
  3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict
  4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

# Conflicting Instructions

- Non-conflicting actions: Let  $T_i$  and  $T_j$  be two different transactions ( $i \neq j$ ), then:
  - $r_i(X); r_j(Y)$  is never a conflict, even if  $X = Y$ .
  - $r_i(X); w_j(Y)$  is not a conflict provided  $X \neq Y$ .
  - $w_i(X); r_j(Y)$  is not a conflict provided  $X \neq Y$ .
  - Similarly,  $w_i(X); w_j(Y)$  is also not a conflict, provided  $X \neq Y$ .

*continued...*

- Three situations of conflicting actions (where we **may not swap** their order)
  - **Two actions of the same transaction.**
    - ▶ e.g.,  $r_i(X);w_i(X)$
  - **Two writes of the same database element** by different transactions.
    - ▶ e.g.,  $w_i(X);w_j(X)$
  - **A read and a write of the same database element** by different transactions.
    - ▶ e.g.,  $r_i(X);w_j(X)$
  
- To summarize, any two actions of different transactions may be swapped unless:
  - They involve the same database element, and
  - At least one of them is a write operation.

# Converting conflict-serializable schedule to a serial schedule

- S:  $r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B);$
- $r_1(A); w_1(A); r_2(A); \underline{w_2(A); r_1(B)}; w_1(B); r_2(B); w_2(B);$
- $r_1(A); w_1(A); \underline{r_2(A); r_1(B)}; w_2(A); w_1(B); r_2(B); w_2(B);$
- $r_1(A); w_1(A); r_1(B); r_2(A); \underline{w_2(A); w_1(B)}; r_2(B); w_2(B);$
- $r_1(A); w_1(A); r_1(B); \underline{r_2(A); w_1(B)}; w_2(A); r_2(B); w_2(B);$
- $r_1(A); w_1(A); r_1(B); w_1(B); r_2(A); w_2(A); r_2(B); w_2(B);$

# conflict-serializable schedule

- **Schedule 3 can be transformed into Schedule 6**, a serial schedule, by series of swaps of non-conflicting instructions..

$T_1$	$T_2$	$T_1$	$T_2$
read (A) write (A)	read (A) write (A)	read (A) write (A) read (B) write (B)	
read (B) write (B)	read (B) write (B)		read (A) write (A) read (B) write (B)

Schedule 3

# conflict-serializable schedule

$T_1$	$T_2$	$T_1$	$T_2$
read (A) write (A)	read (A) write (A)	read (A) write (A) read (B) write (B)	
read (B) write (B)	read (B) write (B)		read (A) write (A) read (B) write (B)

Schedule 3

- **S3:** r1(A), w1(A), r2(A), w2(A), r1(B), w1(B), r2(B), w2(B)
- r1(A),w1(A),r2(A),r1(B),w2(A), w1(B), r2(B), w2(B)
- r1(A),w1(A),r2(A),r1(B), w1(B), w2(A), r2(B), w2(B)
- r1(A),w1(A),r1(B),r2(A),w1(B), w2(A), r2(B), w2(B)
- **S6:** r1(A),w1(A), r1(B), w1(B),r2(A), w2(A), r2(B), w2(B)



# conflict-serializable schedule

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions. Therefore Schedule 3 is conflict serializable.

$T_1$	$T_2$
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

Schedule 3

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	read (A) write (A) read (B) write (B)

Schedule 6

# conflict-serializable schedule

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read ( $Q$ )	write ( $Q$ )
write ( $Q$ )	

- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

# Conflict Serializability & Graphs

- Theorem: **A schedule is conflict serializable iff its precedence graph is acyclic.**
- Theorem: 2PL ensures that the precedence graph will be acyclic!
- Strict 2PL improves on this by avoiding cascading aborts, problems with undoing WW conflicts; i.e., ensuring recoverable schedules.

# Precedence Graphs


- For a schedule  $S$ , involving transactions  $T_1$  and  $T_2$  (among other transactions), we say that  $T_1$  takes precedence over  $T_2$  (written as  $T_1 <_s T_2$ ) if there are actions  $A_1$  of  $T_1$  and  $A_2$  of  $T_2$ , such that:
  - $A_1$  is ahead of  $A_2$  in  $S$ ,
  - Both  $A_1$  and  $A_2$  involve the same database element, and
  - At least one of them is a write operation.

Example of a precedence graph:

Consider a schedule  $S$  which involves three transactions  $T_1$ ,  $T_2$  and  $T_3$ , i.e.,

$S: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$

The precedence graph for this as is shown below



Do not consider  
Same database, same  
transaction, Read – Read,  
which will not conflict

# Test for conflict-serializability

- Construct the precedence graph for S and observe if there are any cycles.

- If yes, then S is not conflict-serializable
- Else, it is a conflict-serializable schedule.

- Example of a cyclic precedence graph:

- Consider the below schedule

$S_1$ :  $r_2(A)$ ;  $r_1(B)$ ;  $w_2(A)$ ;  $r_2(B)$ ;  $r_3(A)$ ;  $w_1(B)$ ;  $w_3(A)$ ;  $w_2(B)$ ;

The precedence graph for this as shown below:




Figure 2

# Precedence Graph- Algorithm

1. Create a node  $T$  in the graph for each participating transaction in the schedule.
2. Draw an edge from  $T_i$  to  $T_j$  in the graph for the conflicting operation.
  - `read_item(X)` and `write_item(X)`
  - `write_item(X)` and `read_item(X)`
  - `write_item(X)` and `write_item(X)`
3. **The Schedule  $S$  is serializable if there is no cycle in the precedence graph.**

# Precedence Graph - Conflict serializable

■ S : r1(x) r1(y) w2(x) w1(x) r2(y)


● x: r1(x) w2(x) w1(x)

● y: r1(y) r2(y)

■ Step 1: Make two nodes corresponding to Transaction  $T_1$  and  $T_2$ .



■ Step 2: Draw edges for the conflicting pair




■ Step 3: Check if there is any cycle formed in the graph. If there is no cycle found, then the schedule is conflict serializable otherwise not

Since the above graph is cyclic, we can conclude that it is **not conflict serializable**

# Example

■ S1:  $r1(x)$   $r3(y)$   $w1(x)$   $w2(y)$   $r3(x)$   $w2(x)$

- x:  $r1(x)$   **$w1(x)$**   $r3(x)$   **$w2(x)$**
- y:  $r3(y)$   $w2(y)$




Since the graph is acyclic, the schedule is conflict serializable.



# Example

- Check whether the given schedule S is conflict serializable or not-

- S :  $R_1(A)$  ,  $R_2(A)$  ,  $R_1(B)$  ,  $R_2(B)$  ,  $R_3(B)$  ,  $W_1(A)$  ,  $W_2(B)$
- A:  $R_1(A)$  ,  $R_2(A)$  ,  $W_1(A)$
- B:  $R_1(B)$  ,  $R_2(B)$  ,  $R_3(B)$  ,  $W_2(B)$




- Clearly, there exists a cycle in the precedence graph.
- Therefore, the given schedule S is not conflict serializable.

# Exercise

Time	Transaction T1	Transaction T2	Transaction T3
t1	Read(X)		
t2			Read(Y)
t3			Read(X)
t4		Read(Y)	
t5		Read(Z)	
t6			Write(Y)
t7		Write(Z)	
t8	Read(Z)		
t9	Write(X)		
t10	Write(Z)		

# Solution



- Since the graph is acyclic, the schedule is conflict serializable.

# View Serializability

- Let  $S$  and  $S'$  be two schedules with the same set of transactions.  $S$  and  $S'$  are **view equivalent** if the following three conditions are met, for each data item  $Q$ ,
  1. **Initial reads:** If in schedule  $S$ , transaction  $T_i$  **reads the initial value of  $Q$** , then in schedule  $S'$  also transaction  $T_i$  must **read the initial value of  $Q$** .
  2. **W-R Conflict:** If in schedule  $S$  transaction  $T_i$  executes **read( $Q$ )**, and that value was produced by transaction  $T_j$  **—write( $Q$ )** (if any), then in schedule  $S'$  also transaction  $T_i$  must read the value of  $Q$  that was produced by the same **write( $Q$ )** operation of transaction  $T_j$ .
  3. **Final write:** The transaction (if any) that performs the **final write( $Q$ )** operation in schedule  $S$  must also perform the **final write( $Q$ )** operation in schedule  $S'$ .

As can be seen, view equivalence is also based purely on **reads** and **writes** alone.

# View Serializability

## ■ Initial reads

- S1: r1(A)
- S2: r1(A)

T1	T2	T1	T2
<hr/>		<hr/>	
R(A)		R(A)	
	W(A)	W(A)	
W(A)			W(A)

## ■ W-R Conflict:

- These two schedule are not view-equivalent:
- S1: w2(A), r3(A)
- S2: w1(A), r3(A)

T1	T2	T3	T1	T2	T3
<hr/>			<hr/>		
W(A)			W(A)		
	W(A)				R(A)
		R(A)		W(A)	

## ■ Final write

- S1: w1(A)
- S2: w1(A)

T1	T2	T1	T2
<hr/>		<hr/>	
R(A)		R(A)	
	W(A)	W(A)	
W(A)			W(A)

# View Serializability

- S1:  $r1(P)$ ,  $w2(P)$
- S2:  $w2(P)$   $r1(P)$
- Initial reads
  - Step1: Satisfied
- W-R Conflict
  - Step2: **Not Satisfied**
- Final write
  - Step3: Satisfied
- Conclusion
  - Not view serializable

Schedule S1:

Time	Transaction T1	Transaction T2
T1	Read (P)	Write (P)
T2		

Schedule S2:

Time	Transaction T1	Transaction T2
T1	Read (P)	Write (P)
T2		

# View Serializability

- S1: w1(P), r2(P), w3(P)
- S2: r2(P), w1(P), w3(P)
- Initial reads
  - Step1: Satisfied
- W-R Conflict
  - Step2: **Not Satisfied**
- Final write
  - Step3: Satisfied
- Conclusion
  - Not view serializable

Schedule S1:

Time	Transaction T1	Transaction T2	Transaction T3
T1	Write (P)	Read (P)	Write (P)
T2			
T3			

Schedule S2:

Time	Transaction T1	Transaction T2	Transaction T3
T1	Write (P)	Read (P)	Write (P)
T2			
T3			

# Exercise

Schedules S1 and S2 are view equivalent

T1:	R(A)	W(A)
T2:	W(A)	
T3:		W(A)

T1:	R(A),W(A)	
T2:		W(A)
T3:		W(A)



# View Serializability Example

S: r2(B) w2(A) r1(A) r3(A) w1(B) w2(B) w3(B)

---

S': r2(B) w2(A) w1(B) r1(A) w2(B) r3(A) w3(B)

# View Serializability Example

## Step 1: Initial reads

→ r2(B) w2(A) r1(A) r3(A) w1(B) w2(B) w3(B)

---

→ r2(B) w2(A) w1(B) r1(A) w2(B) r3(A) w3(B)

# View Serializability Example

## Step 2: W-R Conflict

r2(B) **w2(A)** **r1(A)** **r3(A)** w1(B) w2(B) w3(B)

---

r2(B) **w2(A)** w1(B) **r1(A)** w2(B) **r3(A)** w3(B)

r2(B) w1(B) w2(B) w3(B)

**w2(A)** **r1(A)** **r3(A)**

# View Serializability Example

## Step 3: Final write

r2(B) w2(A) r1(A) r3(A) w1(B) w2(B) w3(B)

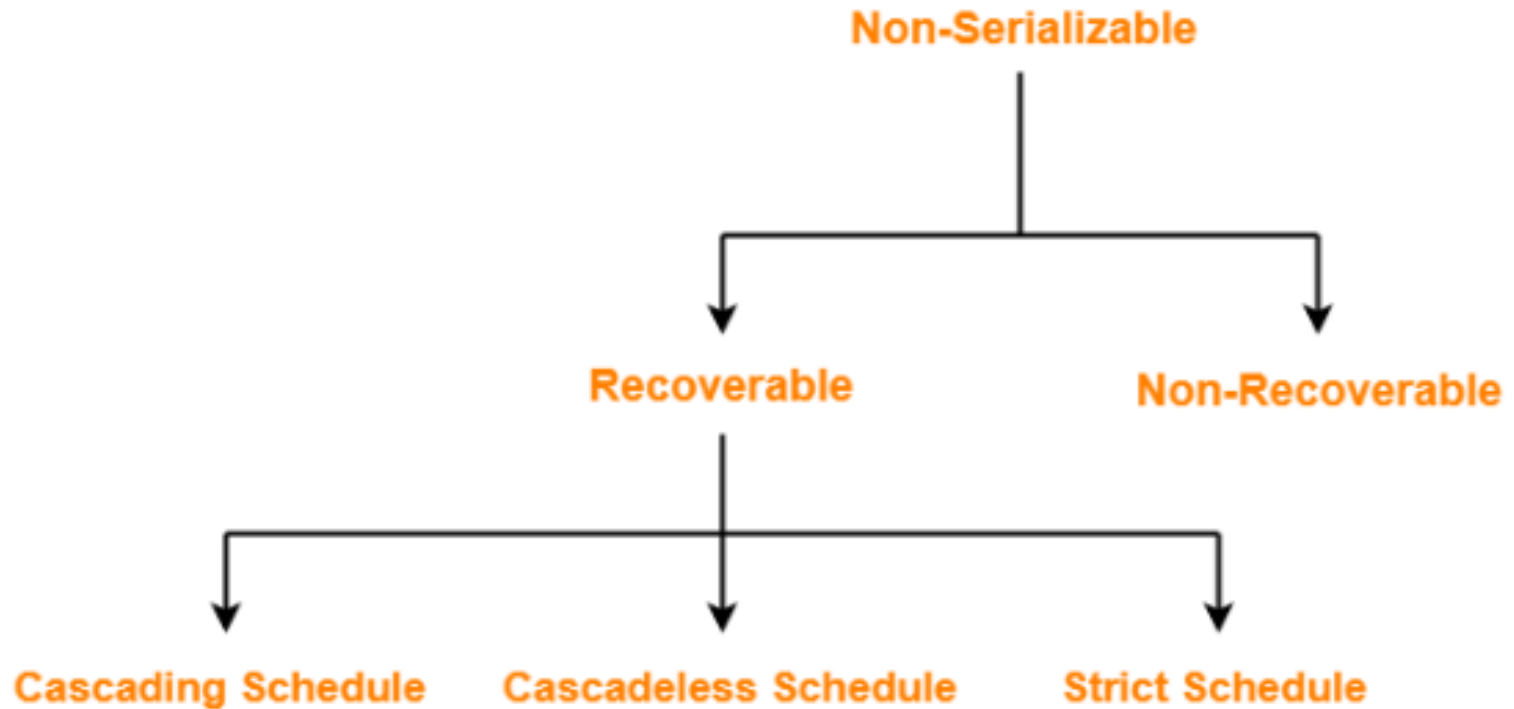
---

r2(B) w2(A) w1(B) r1(A) w2(B) r3(A) w3(B)

# Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction.
- In SQL, a transaction begins implicitly.
- A transaction in SQL ends by:
  - **Commit work** commits current transaction and begins a new one.
  - **Rollback work** causes current transaction to **abort**.

# Recoverable Schedules



# Non-recoverable Schedule.

- If in a schedule,
  - A transaction performs a **dirty read** operation from an uncommitted transaction
  - And **commits before the transaction from which it has read the value**
    - ▶ then such a schedule is known as an **Non-recoverable Schedule**.

# Recoverable Schedules

Need to address the effect of **transaction failures** on concurrently running transactions.

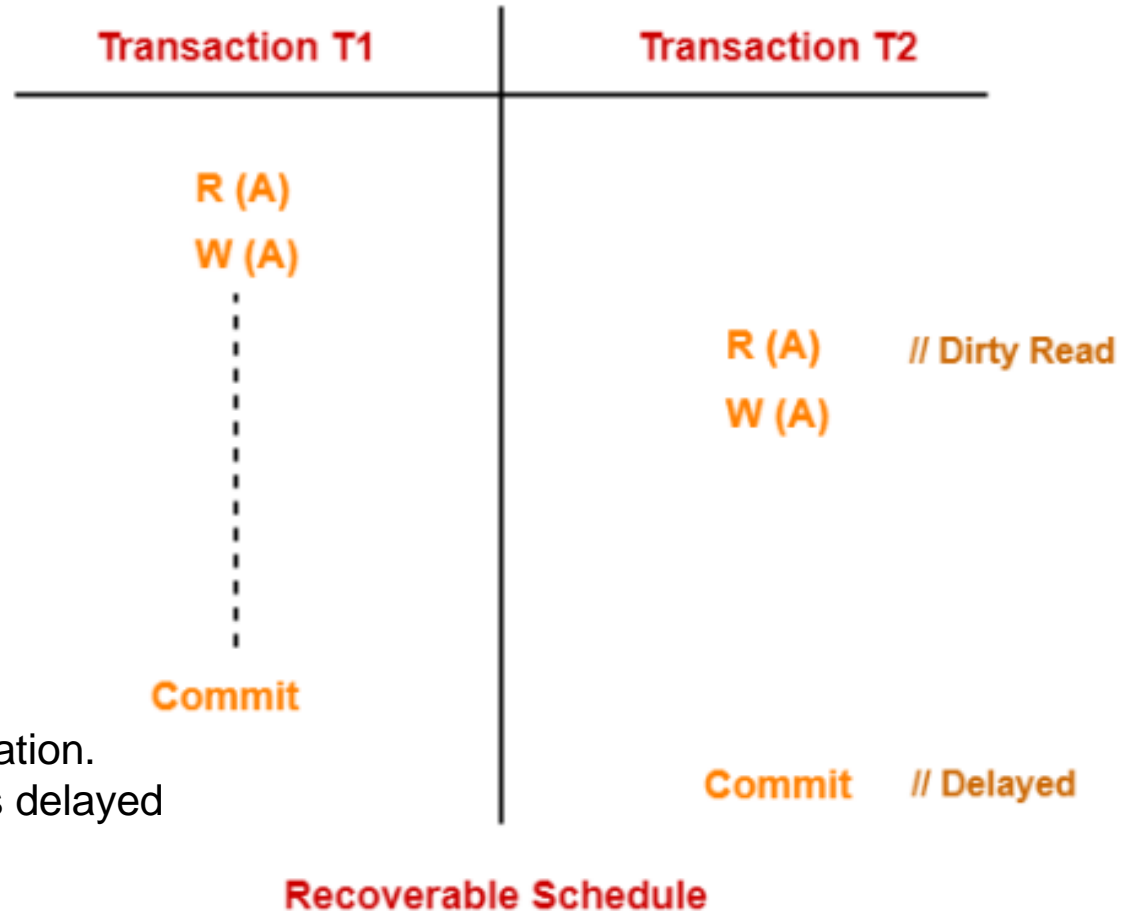
- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- The following schedule (Schedule 11) is not recoverable if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read (A)	
write (A)	
	read (A) // Dirty Read
	commit
read (B)	
Rollback	

- T2 performs a dirty read operation.
- T2 commits before T1.
- T1 fails later and roll backs.
- The value that T2 read now stands to be incorrect.
- T2 can not recover since it has already committed.

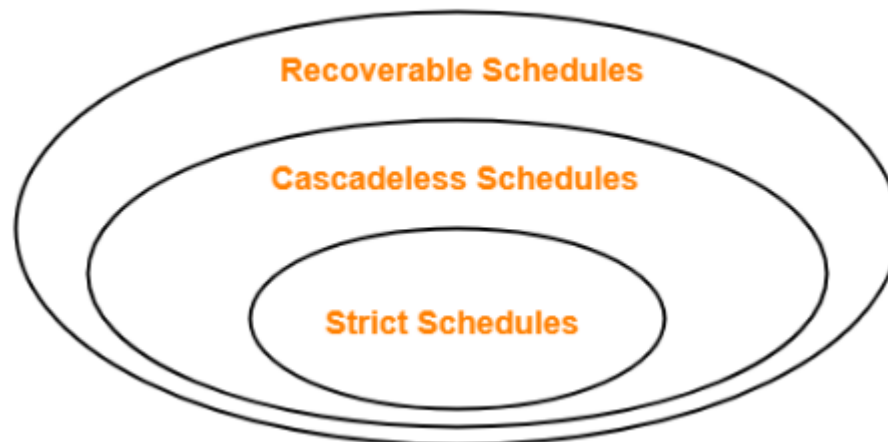
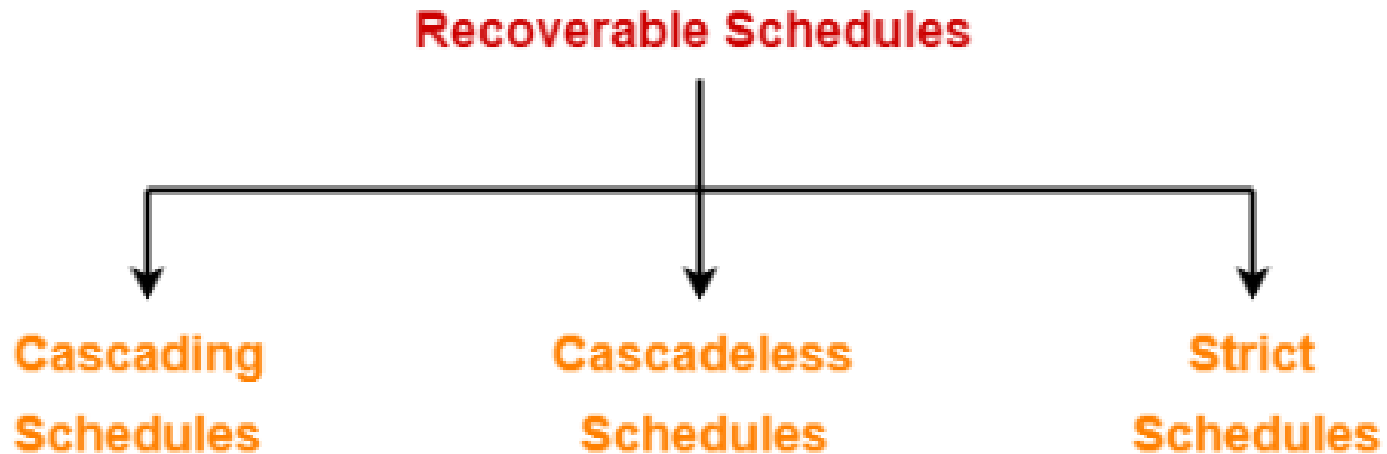


# Recoverable Schedules



- T2 performs a dirty read operation.
- The commit operation of T2 is delayed till T1 commits or roll backs.
- T1 commits later.
- T2 is now allowed to commit.
- In case, T1 would have failed, T2 has a chance to recover by rolling back.

# Types of Recoverable Schedules



# Cascading Rollbacks

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.

- Can lead to the undoing of a significant amount of work

# Cascading Schedules and Rollbacks

T1	T2	T3	T4
R (A)			
W (A)			
	R (A)		
	W (A)		
		R (A)	
		W (A)	
			R (A)
			W (A)
abort / Failure			

Cascading Recoverable Schedule

Transaction T2 depends on transaction T1.

Transaction T3 depends on transaction T2.

Transaction T4 depends on transaction T3.

The failure of transaction T1 causes T2, T3 and T4 to rollback.

Such a rollback is called as a **Cascading Rollback**.

# Cascadeless Schedule

- If in a schedule, a transaction is **not allowed to read a data item** until the last transaction that has written it is committed or aborted, then such a schedule is called as a **Cascadeless Schedule**.

T1	T2	T3
R (A)		
W (A)		
Commit		
	R (A)	
	W (A)	
	Commit	
		R (A)
		W (A)
		Commit


**Cascadeless Schedule**

T1	T2
R (A)	
W (A)	
Commit	
	W (A) // Uncommitted Write

**Cascadeless Schedule**

# Strict Schedule

- If in a schedule, a transaction is **neither allowed to read nor write a data item** until the last transaction that has written it is committed or aborted, then such a schedule is called as a **Strict Schedule**.



## References

- Abraham Silberschatz, Henry F. Korth and S. Sudarshan, *Database System Concepts*, 6th Edition, McGraw-Hill
- <https://www.gatevidyalay.com/serializability-in-dbms-conflict-serializability/>